

From Text to Vectors

Words and Representations

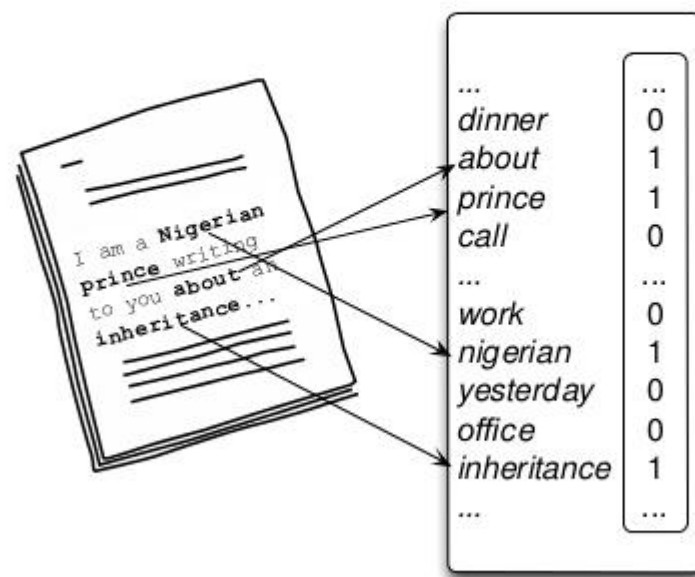
Snorre Ralund, Ph.D Fellow, SoDaS, UCPH

KØBENHAVNS UNIVERSITET



Text to Vector

bag of words



Text to Vector

Document Term Matrix – BoW (Bag of Words)

- Word counts as columns.
- > map word to an index in a matrix.
- Throw out word order.

Text to Vector

Document Term Matrix – BoW (Bag of Words)

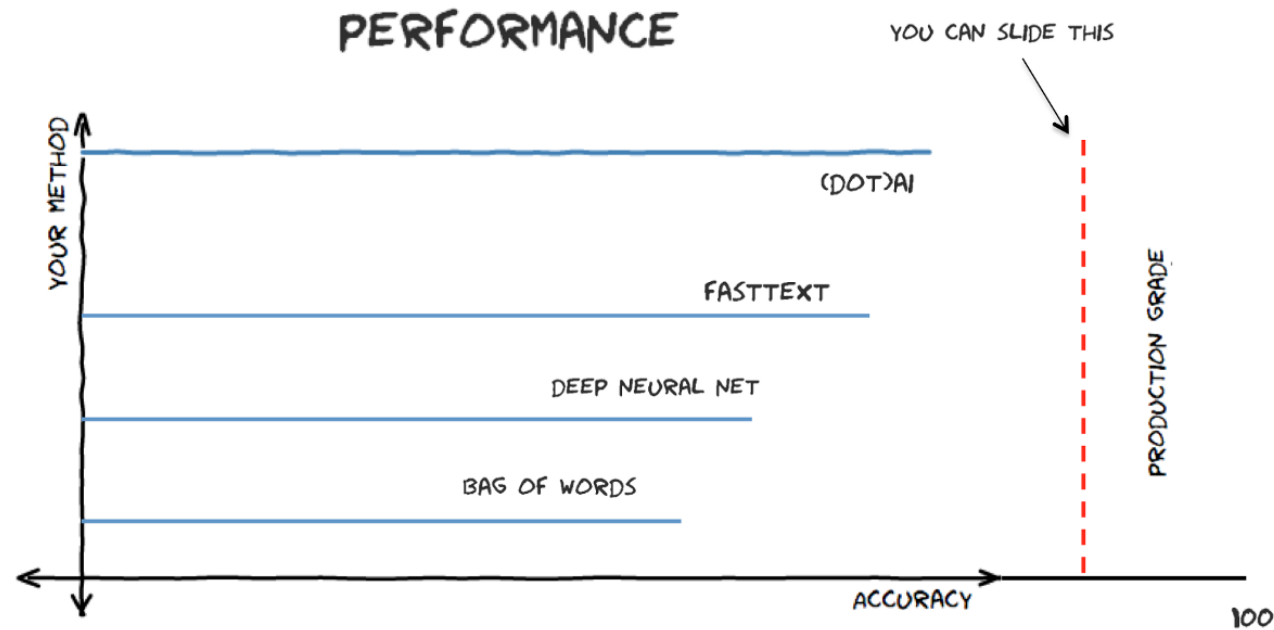
- doc1: *"i really do like bacon"*
- doc2: *"i really do not like bacon"*

Text to Vector

Document Term Matrix – BoW (Bag of Words)

- doc1: *"i really do like bacon"*
- doc2: *"i really do not like bacon"*

	really	do	like	i	bacon	not
doc1	1	1	1	1	1	0
doc2	1	1	1	1	1	1



Baseline models: Naive Bayes, Logistic Regression, K-nearest Neighbor and Support Vector Machines

- **Wang and Manning 2012** *"Baselines and Bigrams: Simple, Good Sentiment and Topic Classification"*:

- State-of-the-art (2012) Topic and Sentiment Classification using only atomized Words as input(BoW) to a linear model.
- No grammar or reasoning.

Bag of Words(2) - Problem with polesemy


Consider the following to documents:

doc1: *"River A/S declared default by the bank."*

doc2: *"When camping my default is by the river bank."*

document	declared	by	default	bank	river	a/s	when	camping	my
doc1	1	1	1	1	1	1	0	0	0
doc2	0	1	1	1	1	0	1	1	1

Bag of Words(2) - Problem with polesemy



Bag of Words (3) - Lack of word orderI

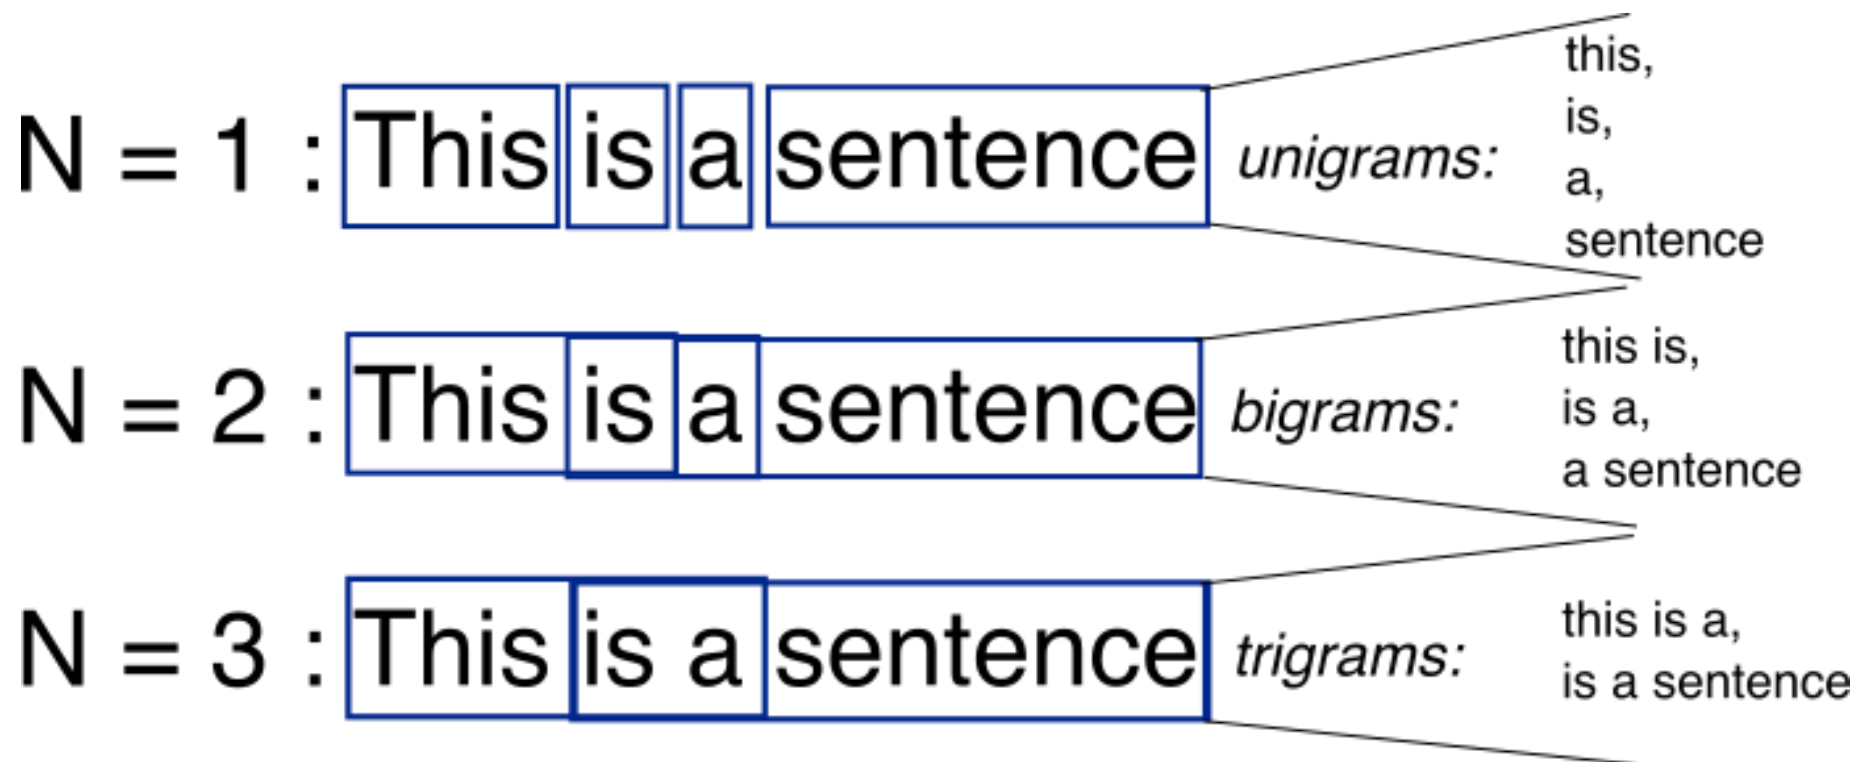
doc1 = 'this was not the best movie'

doc2 = 'was this not the best movie ever?'

Will have very similar representations.



Ngrams to the Rescue



Ngrams to the Rescue(2): Problem with dimensionality

- Quad-grams Quint Grams etc. Generates exponential number of features.

Solution

- Pick only the ngrams using statistical analysis of the word co-occurrences.
- Check out methods for doing so in the Natural Language Processing Toolkit package nltk: ``nltk.collocations`` and the package ``sumgram``





```
"What are the boundaries between-words and meaning?".split('-')  
['What are the boundaries between', 'words and meaning?']
```

Tokenization (1)I

- How we "split" / locate words in text determines the number of dimensions (columns in the Document-Term Matrix)

Issues

- Computational inefficiency
- Parameters are not shared among equivalent words.
 - It makes a difference especially for low N tasks.
 - E.g. run, ran, runs will not share parameters.

Tokenization (2)I

Issues

- Spelling mistakes, or weird uses of punctuation
 - fuzzy matching: ``fuzzymatching package``
- Emoticons: </3 , (:) , :-]
- Multiwords: #no-more-work, New York, Federal Bureau of Finance, word/concept
 - Collocations: ``sumgram package``

Representation (1)

How to encode all relevant information in our tokens?

- lower-casing: DO YOU REALLY WANT TO IGNORE MY ALLCAPS?!?!
 - Our featurespace can potentially double if we don't lowercase.
- Numbers: Infinite combinations of digits
- Filtering to reduce dimensions: Which words to lose?
 - Common or Rare?
 - TF-IDF

$$w_{i,j} = tf_{i,j} \times \log\left(\frac{N}{df_i}\right)$$

$tf_{i,j}$ = number of occurrences of i in j

df_i = number of documents containing i

N = total number of documents

Representation (2)

Grammatical Forms: Do we need grammar?

NO

- Stemming
 - Rulebased: Strips common endings: 'ing','ly','s'
- Lemmatization
 - Lookup in Lexical Ressources(e.g. WordNet): ran --> run, running --> run

Trade-off precision and coverage

Representation (3)

Grammatical Forms: Do we need grammar?

YES


- Use NLP systems for parsing the text

Implementations

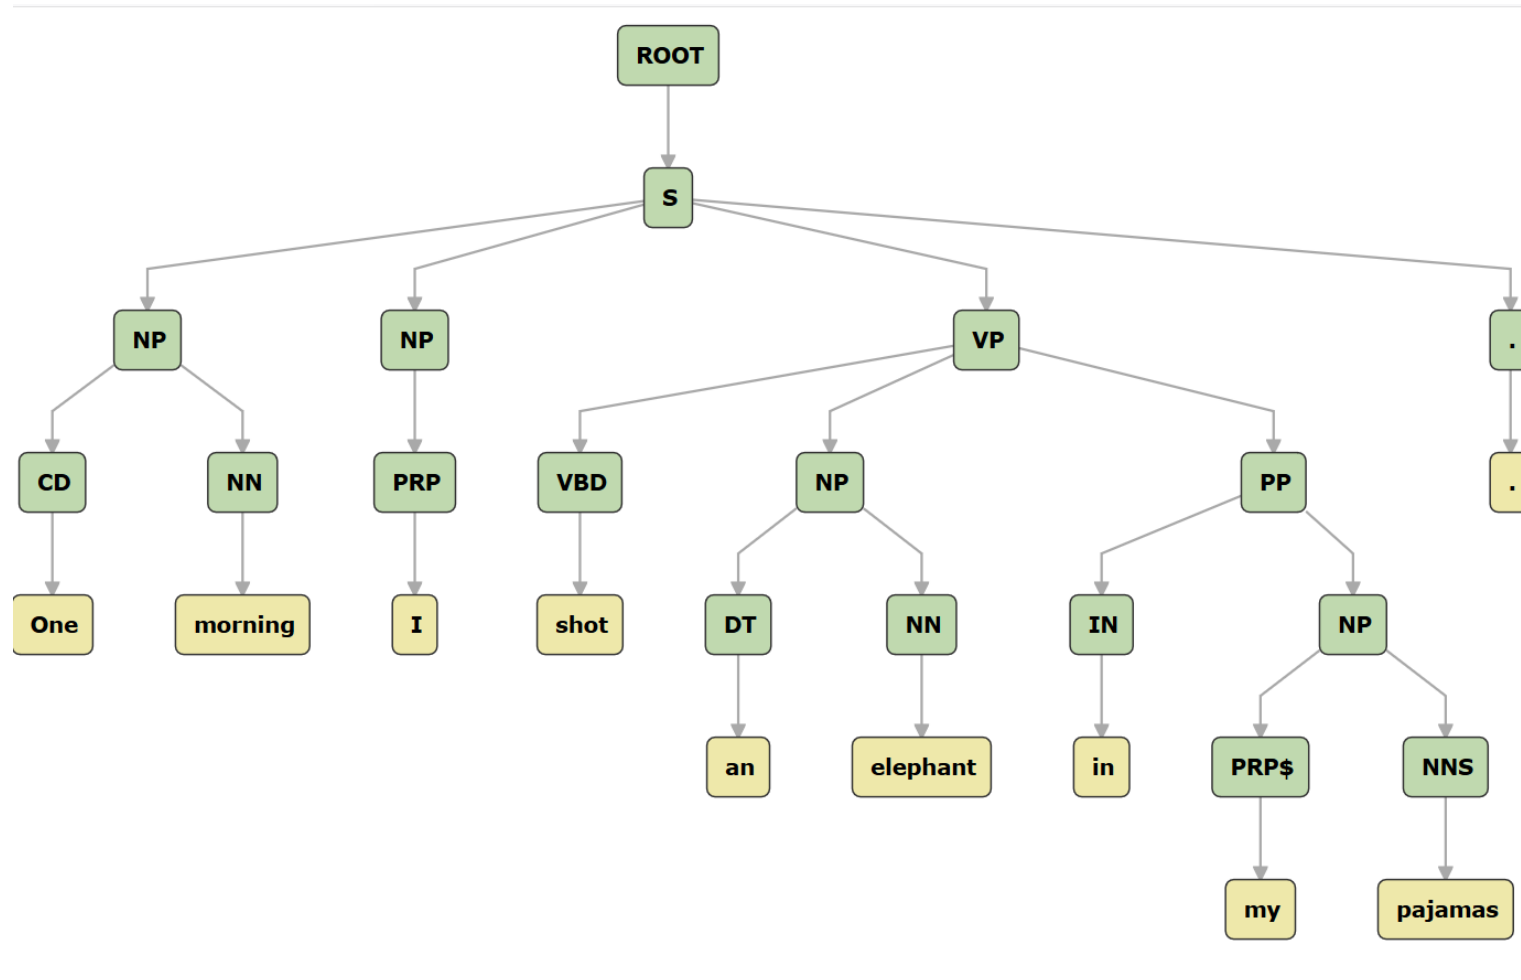
stanfordnlp package: <https://stanfordnlp.github.io/stanfordnlp/>

SpaCy package: <https://spacy.io/>

Representation: Wordsense and Relationship Parsing

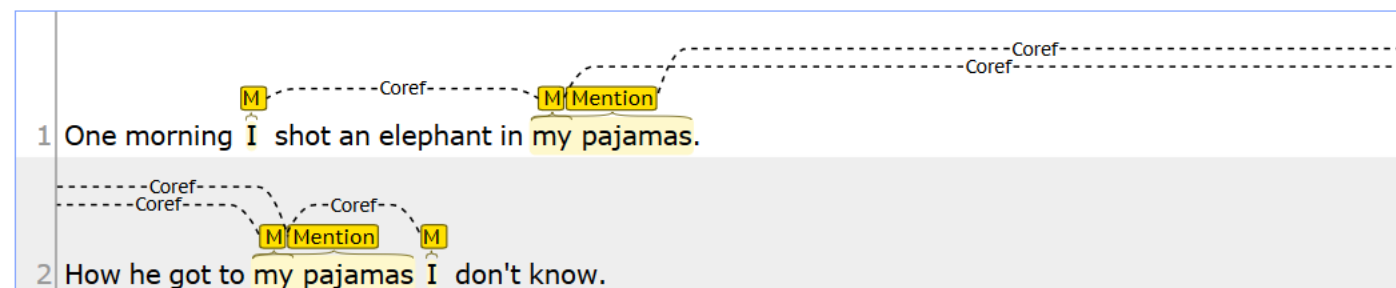


Wordsense and Relationship Parsing

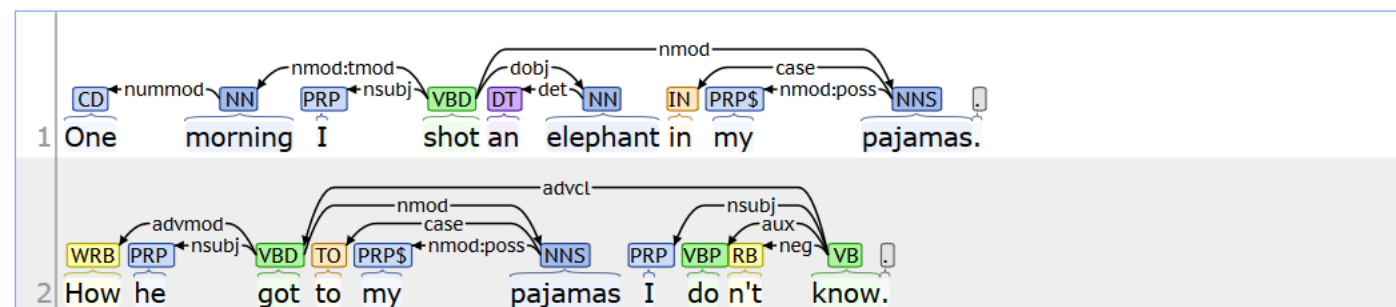


Representation: Wordsense and Relationship Parsing

Coreference:



Basic Dependencies:



Wordsense and Relationship Parsing

- Now we can distinguish between "to run" "a run".
- We have now encoded that "Worst" and "Worse" has the same root lemma.
- Furthermore it can be used to create simple rules about **subject-verb-object** relations (more on this later).
 - Relationships: He (SUBJECT) loves (VERB) her (OBJECT).
 - Simple rule: Unrequitted Love: John (SUBJECT) loves (VERB) Jane (OBJECT) + Absence of reciprocity or Jane (SUBJECT) loves (VERB) Joe (OBJECT)




Tokenization: Choosing a tokenizer

- Build around regular expressions assuming whitespace as separator.
 - e.g. `'[a-z]+|[0-9]+|[\^\sa-z0-9]'`
- Hardcoding complex rules to capture:
 - Emojis. E.g. `=-) <3`
 - Abbreviations e.g. *Ph.D.*
 - Formulas: e.g. $C_{11}H_{12}N_2O_2$ $C_{13}H_{16}N_2O_2$ $C_8H_{11}NO_2$ $C_{28}H_{44}O$ – State-of-the-art



"Klatret^asen(Catch That Girl) is really great movie! It's a 'happy' movie. I watched this movie in 'Puchon International Fantastic Film Festival(PiFan)' on July 12nd, 2003. There is Action + Adventure + Comedy + Thrill + Happy + Romance(cute kids' love Triangle!). You must see this movie. :)"

Tokenization: Choosing a tokenizer



Klatret[©]asen(Catch That Girl) is really great movie! It's a 'happy' movie. I watched this movie in 'Puchon International Fantastic Film Festival(PiFan)' on July 12nd, 2003. There is Action + Adventure + Comedy + Thrill + Happy + Romance(cute kids' love Triangle!). You must see this movie. :)

Tokenization: Choosing a tokenizer



Tokenization: Choosing a tokenizer

For user generated content and social media data use:

- If enough data: `nltk.tokenize.casual.tweet`
- If smaller data: `spacy` or `stanfordnlp` lemmatizer.

For more formal text (e.g. scientific articles) test a few others.

```
import nltk
tweet_tokenizer = nltk.tokenize.casual.TweetTokenizer()
tweet_tokenizer.tokenize('hello I speak emoticon and #hashtag :)')
```

Subword tokenization

Problems

- Dimensionality constrains (many many words) and unseen words in relation a fixed Vocabulary
- Words can be compositional: e.g.
"Speciallægepraksisplanlægningsstabiliseringsperiode"
- Grammer: You want to be able to learn similarities between 'worse' and 'worst'.

Subword tokenization


Problems

- Dimensionality constrains (many many words) and unseen words in relation a fixed Vocabulary
- Words can be compositional: e.g.
"Speciallægepraksisplanlægningsstabiliseringsperiode"

Subword tokenization

Solutions

- Characterbased LSTM
 - Can work, but has long training times. E.g. "<https://openai.com/blog/unsupervised-sentiment-neuron/>
 - Maybe the "embedding" of a character has to serve too "heterogenous" purposes / i.e. complex combinatorial transformations.
 - Remember the Idea of a RNN.



Subword tokenization

Solutions

- Subword tokenization
 - Preprocess by extracting subwords as highly co-occurring characters sequences.
- BytePairEncoding: Seinrich et. al 2016 "Neural Machine Translation of Rare Words with Subword Units".
 - "is a simple data compression technique that iteratively replaces the most frequent pair of bytes in a sequence with a single unused byte".

Example: Subword tokenization of a **RAP** song



Example: Subword tokenization of a rap song

"I said a hip, hop, the hippie, the hippie, to the hip hip-hop, and you don't stop".

- **"op"**, **"ip"** and **"he"** most frequent.

→ substitute for byte 0 1 and 2

"I said a h0, h1, t2 h0pie, t2 h0pie, to t2 h0 h0-h1, and you don't st1"

- Now **"0p"** most frequent.
- "I said a h0, h1, t2 h3ie, t2 h3ie, to t2 h0 h0-h1, and you don't st1". **"3i"** most frequent.