

Transfer Learning in NLP

Self-supervised learning tasks and model capacity in NLP

Snorre Ralund, Ph.D Fellow, University of Copenhagen, SoDaS

KØBENHAVNS UNIVERSITET



Transfer Learning in NLP

Self-supervised learning tasks and model capacity in NLP

Snorre Ralund, Ph.D Fellow, University of Copenhagen, SoDaS



KØBENHAVNS UNIVERSITET



Self-Supervised Learning

- Extract and use the metadata and relevant context as supervisory signal.

Self-Supervised Learning

- Extract and use the metadata and relevant context as supervisory signal.
- Self-referential prediction tasks.

E.g. Leverage large scale user generated tags

- Hashtags for expressing topic and summarization.
 - #sarcasm #irony
 - #topic
- Emojies to explicate emotional intention - DeepMoji

Self-Supervised Learning

- Extract and use the metadata and relevant context as supervisory signal.
- Self-referential prediction tasks.

E.g. Leverage large scale user generated tags

- Hashtags for expressing topic and summarization.
 - #sarcasm #irony
 - #topic
- Emojies to explicate emotional intention - DeepMoji

Other User generated tags

- Subreddits - including reactions
- Keywords in Scientific Articles
- Tags in stackoverflow

Self-Supervised Learning

Language models: Supervisory signal from raw text

- Predict next word / Char given previous word
- Predict word given context. (Cooperation and interactions)
 - Word2Vec - Context window.
 - BERT – Removing random words from larger (32) Context Window
- Predict word given previous context + Reverse (ELMO)
- Predict next / previous sentence
 - E.g. SkipThoughts, BERT.
- Denoising better than Language models (Raffel et. al 2019)


External Reference

- View the Director of One of the Major Transfer Learning Hubs, introduction to transfer learning.
 - https://www.youtube.com/watch?v=0T_Qr4qBrqc




Reusing models

- Large Expensive Models can be re-used.
- Sharing via Hubs:
 - Transformers package
 - TFHUB
 - PYTORCH HUB
- Framework agnostic.




Model Capabilities: Input - Encode – Decode - Output



Linear regression equation
(without error)

$$\hat{Y} = bX + a$$

predicted values of \hat{Y} slope = rate of increase/decrease of \hat{Y} for each unit increase in X Y-intercept = level of \hat{Y} when X is 0.



Model Capabilities

- Encode information from input into a Vector (or network of vectors)
 - Syntax, semantics, topical information, facts etc.
 - E.g.
Dimensions: [Noun, Active, Animal]
 - Mouse -> [0.2,-0.3,1]
 - Cat -> [0.2, 0.3, 1]
 - Catch -> [-0.2,0.3,-0.5]
- Decode information
 - Process the encoded information to produce output.

```
glove_200.most_similar(positive=['police','black'],negative=['white'])[0:5]
```

```
[('cops', 0.7516576051712036),  
 ('officers', 0.6661906838417053),  
 ('arrested', 0.6204742193222046),  
 ('suspect', 0.6187559366226196),  
 ('cop', 0.6156525015830994)]
```

```
glove_200.most_similar(positive=['police','white'],negative=['black'])[0:5]
```

```
[('cops', 0.7516659498214722),  
 ('officers', 0.7105646133422852),  
 ('authorities', 0.6782428026199341),  
 ('arrest', 0.6773560047149658),  
 ('officials', 0.662535548210144)]
```

So which models have what capabilities

- Word2Vec, FastText only simple attenuation.
- 3 layers, and embeddings are averaged.
 - Learns linear information (similar to a BOW)




Figure 1: Model architecture of `fastText` for a sentence with N ngram features x_1, \dots, x_N . The features are embedded and averaged to form the hidden variable.

Simple Linear Transformation – Attenuation

Dimensions to be encoded



"hidden 0"



"I"

"hidden 1"



"am"

"hidden 2"

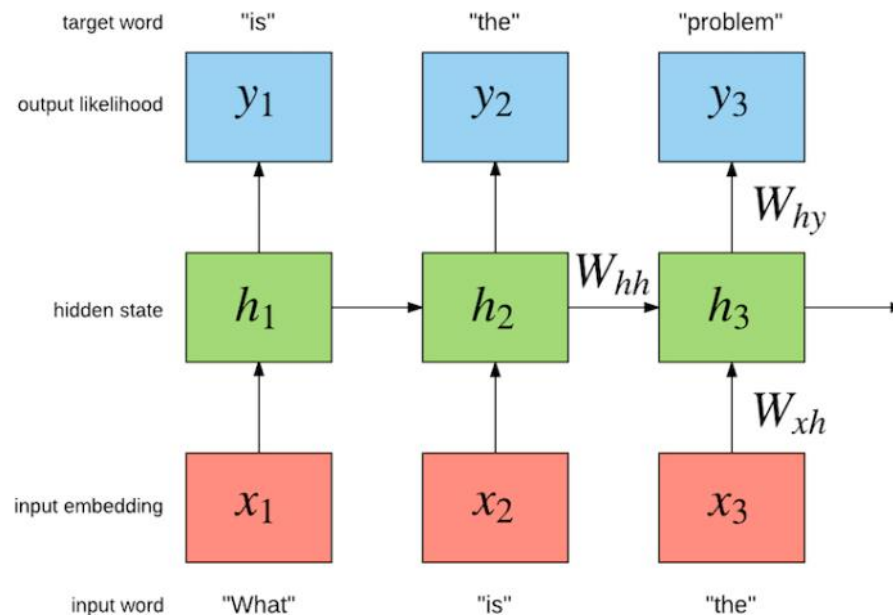


"happy"

So which models have what capabilities

ELMo - Embeddings from Language Models. "Deep Contextualized word representations"

- Forward reading updating hidden states.
- Sequence of hidden states.



Simple Linear Transformation - Interaction

Dimensions to be encoded



"hidden 0"




"hidden 1"




"hidden 2"




Non-linearities: Re-routing and Memory



Non-linearities: Re-routing and Memory



Non-linear transformation: Epistatic and Multiple Interactions – Threshold



A by itself = NO
Activation


B by itself = NO
Activation

A+B = Activation

So which models have what capabilities

ELMo - Embeddings from Language Models. "Deep Contextualized word representations"

- Bi-directional: Forward and Backward.
- But only Sequential information



So which models have what capabilities

Bi-Directional models – Transformer Models

- BERT, GPT2, Electra ..
 - Allows word embeddings to interact by transforming each other.
 - Multiple transformation of each embedding is allow to interact multiple times, to allow for even more complex interactions.
 - See for a more detailed discussion: <https://www.youtube.com/watch?v=ycXWAtm22-w>

So which models have what capabilities

- Memory – Short and Long term.
 - Models struggle to incorporate long term knowledge in longer texts ([Dailuk et al. 2017](#))
 - Open research question: Rae et. al 2019: "COMPRESSIVE TRANSFORMERS FOR LONG-RANGE SEQUENCE MODELLING"
 - BERT has a fixed context window.
 - XLNET, Transformer XL, Compressive Transformer are working in this direction.

External Reference 2


- View technical exposition of the history of language models
 - <https://www.youtube.com/watch?v=ycXWAtm22-w>

NNML


- A Neural Probabilistic Language Model

$$L = \frac{1}{T} \sum_i \log P(w_i | w_{i-1}, \dots, w_{i-m+1}; \theta) + R(\theta)$$

$$P(w_i | w_{i-1}, \dots, w_{i-m+1}; \theta) = \frac{e^{h_{w_i}}}{\sum_c e^{h_c}}$$



Transfer Learning: Adaptation



Transfer Learning: Adaptation

- Match Tokenization Scheme.


```
# Glove preprocessing (rewritten for python) for following ruby script:
#https://nlp.stanford.edu/projects/glove/preprocess-twitter.rb
import re
RE_URL = r'(?:(https?://www\.)?(?:[a-zA-Z]|[0-9]|[_-@.&+]|[*%\{\}])|(?:[0-9a-zA-Z-]{0-9a-zA-Z-})+)'
RE_EMAIL = r'(?!\b[a-zA-Z0-9_-]+(?:@)[a-zA-Z0-9-]+\.[a-zA-Z0-9-]+(?:\.[a-zA-Z0-9-]+)?\s?)'
RE_USER = r'(?!(?!\w+)(?:[-|_|!])?'
RE_HEART = r'(?:</?3+)' # including broken heart.
RE_HEART = r'<3'
# Construct emoticons
eyes = r"[8;=;]"
noses = r"[\^`~]?"
mouth = r'([dD])'
mouth_right = r'([l])'
smiles = [eyes+noses+mouth,mouth_right+noses+eyes]
SMILE_RE = '|'.join(smiles)
LOL_RE = eyes+noses+'p'
frowns = [mouth+noses+eyes,
          eyes+noses+mouth_right]
SAD_RE = '|'.join(frowns)
neutral_mouth = r'([v/!])'
NEUTRAL_RE = eyes+noses+neutral_mouth
# Make sure it is not a hashtag.
ALLCAPS_RE = r'(?!(?!(?!\b(?:[A-Z0-9_]*)?[A-Z0-9]{2,}|(?:[A-Z0-9_]*)?(?:\b|(?![.!?]))'
test = '#ASHDSH ASDRASLD90234 #FASH'
HASHTAG_ALLCAPS_RE = r'(?!(?!(?!\b(?:[A-Z0-9_]*)?[A-Z]{2,}|(?:[A-Z0-9_]*)?(?:\b|(?![.!?]))'
HASHTAG_RE = r'(?!(?!(a-zA-Z0-9_]+)'
regex_replace = [(RE_URL, '<URL>'),
                 (RE_URL2, '<URL>'),
                 (RE_USER, "<USER>"),
                 (RE_EMAIL, "<EMAIL>"),
                 ## ALLCAPS
                 (ALLCAPS_RE, r'<ALLCAPS> \1'),
                 #emoticons
                 (SMILE_RE, '<SMILE>'),
                 (LOL_RE, '<LOLFACE>'),
                 (SAD_RE, '<SADFACE>'),
                 (NEUTRAL_RE, "<NEUTRALFACE>"),
                 (RE_HEART, '<HEART>'),
                 ## HASHTAGS
```

Transfer Learning: Adaptation

- Match Tokenization Scheme.
- Streamlining of this is developed around the package tokenizers:
 - <https://github.com/huggingface/tokenizers>

Transfer Learning: Adaptation

- Match Tokenization Scheme
- **Finetuning** of ***Language Model*** for domain adaptation. (Howard and Ruder 2017)



Transfer Learning: Adaptation

Finetuning for Classification

- Transfer techniques
 - Still an open research question, however.

Transfer Learning: Adaptation

Finetuning for Classification

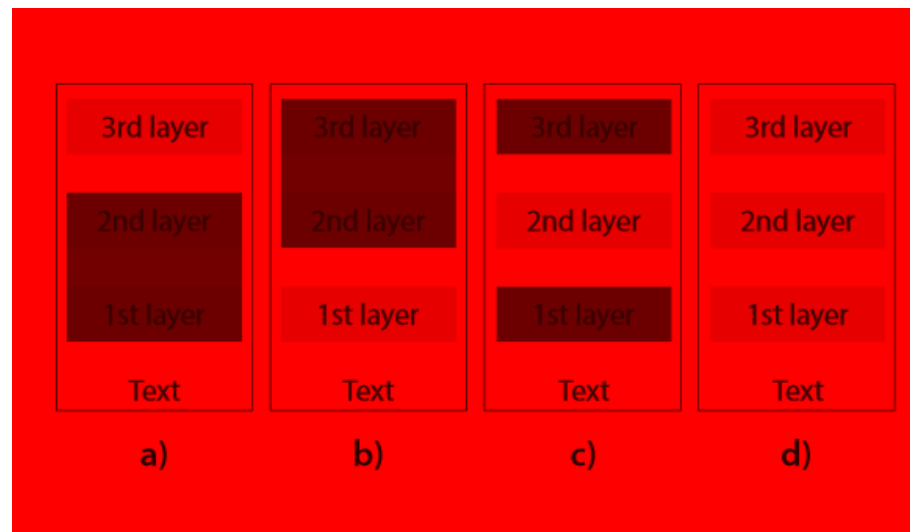
Transfer techniques

- How to Decode?
 - Use word embeddings directly for word level predictions:
 - E.g. parsing.
 - Summarize Embeddings to form Sentence level representation.
- Classifier heads: Simple linear Softmax, or LSTM.

Transfer Learning: Adaptation

Transfer Techniques

- How To Train?
 - Avoid overfitting: large model adopt to small data -- memorization
 - Discriminative learning rates (earlier layers have more fundamental information → lower learning rate → reduce updates.)
 - Gradual unfreezing of layers (Howard and Ruder 2017), Chain-thaw (Felbo et al. 2017)



Transfer Learning: Adaptation

Transfer Techniques

- How To Train?
 - Avoid overfitting: large model adopt to small data -- memorization
 - Discriminative learning rates (earlier layers have more fundamental information → lower learning rate → reduce updates.)
 - Gradual unfreezing of layers (Howard and Ruder 2017), Chain-thaw (Felbo et al. 2017)
 - Auxillary loss function (keeping the Language model objective) (Chronopoulou et. al.2019)

